

MIT 8301 Virtual Laboratory

German Credit Risk Classification

Student	Basil Oforbuike Emeokoro
Matric Number	2025/A/MIT/0111
Student ID	301806653
Email	b.emeokoro6653@miva.edu.ng
Programme	Master of Information Technology
Specialisation	Artificial Intelligence & Machine Learning
Course	MIT 8301 - Machine Learning and AI Fundamentals
Academic Session	2026/2027

Cover Page and Student Details

German Credit Risk Classification

Student: Basil Oforbuike Emeokoro

Matric Number: 2025/A/MIT/0111

Student ID: 301806653

Email: b.emeokoro6653@miva.edu.ng

Programme: Master of Information Technology

Specialisation: Artificial Intelligence & Machine Learning

Course: MIT 8301 - Machine Learning and AI Fundamentals

Academic Session: 2026/2027

Executive Summary

This assessment develops a reproducible classifier for bad credit risk. The supplied file contained 1,000 applicants and nine predictors but omitted the label. An authentic target-inclusive edition was reconciled with a 100% exact row-level predictor match and zero mismatched cells before labels were appended. The verified target contains 700 good and 300 bad applicants. Four models were trained and evaluated with leakage-safe preprocessing. Tuned Logistic Regression was selected under a predeclared performance-and-governance rule.

Problem Statement and Objectives

The objective is to identify risky applicants ($\text{credit_risk}=1$) while recognising the asymmetric institutional and fairness costs of false negatives and false positives. The work covers EDA, missing values, outliers, encoding, scaling, scratch Logistic Regression, tuned model comparison, transparent feature interpretation, and responsible-use limitations.

Dataset Description, Provenance, and Target Restoration

The supplied export had 1,000 rows, an index column, and nine predictors. The target source was the Hugging Face mirror of Kaggle's German Credit Risk - With Target edition. After removing the export index and normalising only whitespace/case, every predictor cell matched. No fuzzy or positional merge was used. good was encoded 0 and bad 1. SHA-256 checksums and per-feature match percentages are recorded in the provenance report.

Exploratory Data Analysis

There are no duplicated complete rows. Age has 53 unique values, credit amount 921, and duration 33. The target is moderately imbalanced (70% good, 30% bad). Histograms, robust-scaled boxplots, categorical distributions, a missingness chart, a correlation heatmap, target comparisons, and observed risk by purpose appear in the figures appendix. These are descriptive associations, not causal findings.

Missing Values and Outliers

Savings status has 183 missing values; checking status has 394. A dedicated Unknown_or_No_Account category preserves potentially informative absence and is learned inside training folds. IQR review found {'Age': 23, 'Credit amount': 72, 'Duration': 70}. Values are plausible, so they are not deleted. Training-fitted 1.5-IQR clipping limits leverage, followed by robust scaling. This avoids data leakage.

Encoding, Scaling, and Split

The data were split into 800 training and 200 untouched test rows with stratification and random state 42. Age, amount, and duration use median imputation, IQR clipping, and RobustScaler. Nominal fields, including Job, use constant imputation and one-hot encoding with unknown-category tolerance. All transformations are fitted only on training data or within cross-validation folds.

Training, Validation and Test Strategy

The dataset was partitioned into an 80% training set (800 observations) and a 20% independent test set (200 observations) using a stratified train-test split with `random_state=42`. The training set contains 560 good and 240 bad applicants, and the test set contains 140 good and 60 bad applicants; both therefore preserve the original 70% good / 30% bad class distribution.

Hyperparameter optimisation was performed exclusively on the training data using Stratified 5-Fold Cross-Validation within GridSearchCV. During this process, the 800 training observations were repeatedly divided into internal training and validation folds while preserving class proportions. Every training observation serves in validation across the rotating folds. After optimal hyperparameters were identified, GridSearchCV automatically refitted each model on the complete 800-row training set before a single evaluation on the untouched 200-row test set.

No separate validation dataset was required because cross-validation already creates internal validation folds. Although 60/20/20 and 70/15/15 partitions are also established approaches, this project intentionally adopted 80% train + 20% test + stratified 5-fold cross-validation because the dataset contains only 1,000 observations, cross-validation provides statistically stronger hyperparameter evaluation, and more observations remain available for model learning. This aligns with current scikit-learn and production machine-learning practice.

Missing-value imputation, categorical encoding, numerical scaling, IQR outlier clipping, and all feature transformations are implemented inside scikit-learn Pipeline and ColumnTransformer objects. They are fitted within training folds and subsequently applied to validation or test rows. GridSearchCV received only `X_train` and `y_train`; the decision threshold remained fixed at 0.5, so the test set influenced neither model selection, hyperparameter tuning, preprocessing, nor threshold optimisation.

Logistic Regression from Scratch

The implementation provides a stable clipped sigmoid, explicit bias, binary cross-entropy, vectorised gradients, L2 regularisation, convergence tolerance, loss history, probability estimates, and input/fitted-state validation. It converged=True in 4233 iterations with final loss 0.50469. Its test performance is compared directly with scikit-learn.

Model Training and Hyperparameter Tuning

Training used stratified five-fold GridSearchCV and ROC-AUC as the primary score. The test set was never used for tuning. Tuning controls regularisation, margin/kernel complexity, class weighting, and Naive Bayes variance smoothing, improving generalisation without test leakage.

- ****Tuned Logistic Regression:**** CV ROC-AUC 0.756; {'model__C': 0.1, 'model__class_weight': None, 'model__penalty': 'l2', 'model__solver': 'liblinear'}
- ****Tuned SVM:**** CV ROC-AUC 0.765; {'model__C': 1, 'model__class_weight': 'balanced', 'model__gamma': 0.1, 'model__kernel': 'rbf'}
- ****Tuned Gaussian Naive Bayes:**** CV ROC-AUC 0.666; {'model__var_smoothing': 1e-12}

Test Evaluation and Comparison

Model	Accuracy	Precision	Recall	F1	ROC-AUC
Tuned SVM	0.710	0.511	0.750	0.608	0.786
Tuned Logistic Regression	0.730	0.594	0.317	0.413	0.766
Logistic Regression from Scratch	0.720	0.548	0.383	0.451	0.756
Tuned Gaussian Naive Bayes	0.675	0.472	0.700	0.564	0.702

A false negative classifies a risky applicant as good and may expose the institution to loss. A false positive classifies a good applicant as risky and may cause unfair rejection or harsher terms. Accuracy alone is therefore insufficient; recall, F1, ROC-AUC, thresholds, interpretability, and operating costs matter.

Best-Model Justification

The predeclared rule prefers tuned Logistic Regression when its ROC-AUC is within 0.03 of the highest model, otherwise the highest-AUC model is selected. The result is Tuned Logistic Regression. This accepts a small discrimination trade-off for transparent coefficients, simpler deployment, auditability, and regulatory suitability. Operational threshold tuning could improve risky-class recall and must be set from validated costs, not the test set.

Feature Importance and Interpretation

The largest absolute transformed Logistic Regression coefficients are:

- cat__Checking account_Unknown_or_No_Account: -0.879
- num__Duration: +0.514
- cat__Checking account_little: +0.471
- cat__Saving accounts_Unknown_or_No_Account: -0.392
- cat__Purpose_radio/tv: -0.360
- cat__Housing_own: -0.343
- cat__Sex_male: -0.330
- cat__Saving accounts_little: +0.313
- cat__Purpose_education: +0.307
- cat__Saving accounts_rich: -0.283

Positive values increase estimated bad-credit log-odds and negative values decrease them, conditional on the encoding and scaling. Coefficients express association, not causation. One-hot terms are interpreted relative to the omitted all-zero representation.

Ethics, Fairness, and Responsible Use

Sex is sensitive and age can create protected-class concerns. Historical labels may encode past discrimination and other fields may serve as proxies. Supplementary metrics by sex are reported but small test subgroups cannot establish fairness. The model requires human review, reasons for adverse decisions, appeal routes, privacy controls, drift monitoring, periodic bias audits, and independent validation. Prediction is not causal judgement.

Limitations

The dataset is small, historical, geographically specific, and omits important affordability and contemporary lending variables. Labels have no event-time detail; calibration and temporal stability are not established. The simplified nine-feature representation loses information from the original 20-feature UCI data. Reported test estimates have sampling uncertainty.

Lessons Learned and Engineering Reflection

Principal engineering challenge

The most consequential challenge was not model selection but the absence of a target variable in the supplied German Credit CSV. Supervised classification requires authentic observed outcomes: without a good/bad credit label, there is no valid response variable from which a classifier can learn or against which it can be evaluated. Creating synthetic labels, inferring them from predictor thresholds, clustering applicants and treating the clusters as outcomes, or deriving labels from the same features used for prediction would have introduced circular reasoning and invalidated the experiment. The appropriate engineering decision was therefore to pause supervised modelling until label provenance could be established. Preserving dataset integrity was more important than forcing a pipeline to produce metrics.

Engineering solution and provenance control

The resolution began with an investigation of the supplied file's schema and the provenance of the simplified German Credit representation. An authentic target-inclusive edition with the same nine predictors was identified. The reconciliation process removed export-only identifiers, standardised column names, and normalised only harmless representation differences such as surrounding whitespace and demonstrably equivalent category case. It then compared every predictor, row by row and feature by feature, before accepting any label.

The verification produced a 100% match for every predictor and zero mismatched cells across 1,000 observations. SHA-256 checksums were recorded for both raw inputs, duplicate diagnostics and class counts were produced, and the procedure was documented in a machine-readable and human-readable validation report. Only after these checks passed was the authentic target appended and mapped to `credit_risk`, with 1 representing bad/risky credit and 0 representing good credit. This retained the original alignment between each applicant and the corresponding outcome.

This experience illustrates why data provenance is an operational requirement rather than an administrative detail. In a real machine-learning system, an apparently small join, row-order change, or unverified label source can silently corrupt the relationship between features and outcomes while leaving the code executable. Checksums, explicit reconciliation rules, fail-fast assertions, and durable validation reports make those risks observable and auditable.

Machine-learning lessons

Data quality. Model quality depends first on the validity, alignment, and meaning of the data. A sophisticated estimator cannot compensate for an unauthenticated target or an incorrect feature-outcome join. The project also reinforced that reproducible preprocessing is more valuable than pursuing a marginal improvement in headline accuracy: the same inputs should pass through the same documented, testable transformations whenever the analysis is repeated.

Preventing leakage. Leakage was prevented by separating training and test data before any learned transformation entered the model workflow. Missing-value imputation, one-hot encoding, robust scaling, and IQR clipping were placed inside Pipeline and ColumnTransformer objects, causing each operation to learn its parameters from the relevant training fold only. The independent test set remained untouched until final evaluation. This matters because leakage can make a weak model appear deceptively successful by allowing information about held-out observations to influence preprocessing, tuning, or selection.

Stratification. The original target distribution contains 70% good credit and 30% bad credit. Stratified sampling preserved this balance exactly in the 800-row training set and 200-row test set. The resulting 560/240 and 140/60 class counts reduced the risk that an accidental shift in class proportions would distort comparison or make risky-class recall less reliable.

Cross-validation. Stratified 5-Fold Cross-Validation was selected instead of reserving a separate validation dataset. With only 1,000 observations, rotating validation folds make better use of the 800 training cases,

provide more stable hyperparameter evidence than one small holdout, and preserve class proportions in every fold. GridSearchCV evaluated configurations only within those training folds and then refitted the selected configuration on the complete training set, supporting a stronger estimate of generalisation before the single test evaluation.

Interpretability. The tuned SVM achieved the highest test ROC-AUC, but tuned Logistic Regression was selected under the predeclared rule because its ROC-AUC was within 0.03 of the leader and its coefficients provide clearer decision logic. In credit-risk settings, engineering quality includes the ability to explain influential factors, reproduce a decision path, and support review by governance teams, auditors, and regulators. Predictive performance remains important, but a small gain from a less transparent model may not outweigh the operational value of understandable and contestable outputs.

Responsible machine-learning reflection

Responsible credit-risk modelling requires controls tied to the actual decision context. Sex and age can expose applicants to disparate outcomes, and apparently neutral variables may act as proxies for protected characteristics. Transparency therefore includes documenting the target, transformations, evaluation population, decision threshold, subgroup diagnostics, and known limitations. Explainability must be usable by reviewers and affected applicants rather than treated as a decorative chart. Reproducible code, versioned data checks, tests, model documentation, approval controls, drift monitoring, and appeal procedures form part of governance. Most importantly, a prediction should support - not replace - professional judgement in a high-impact lending decision; adverse outcomes require human oversight and a meaningful route for review.

Professional reflection

I found that this assignment strengthened my understanding of machine learning as an end-to-end engineering discipline rather than an isolated model-fitting exercise. Restoring the verified target required data engineering and provenance reasoning; the leakage-safe pipeline required careful separation of fitting and transformation; the scratch implementation connected mathematical concepts to tested software; and the tuning and evaluation stages required disciplined validation rather than selective reporting. Packaging the workflow with tests, reproducibility files, generated evidence, documentation, and Git history also demonstrated how software-engineering practices make analytical work easier to inspect, repeat, and maintain. The responsible-AI analysis further clarified that technical validity and accountable use must be designed together in credit-risk applications.

Portfolio value

As a portfolio project, the completed repository demonstrates practical competence in data reconciliation and preprocessing, supervised classification, first-principles model implementation, hyperparameter optimisation, leakage-safe evaluation, explainable modelling, automated testing, reproducible research, report generation, and Git-based project management. The value lies in the traceable workflow from raw data to verified evidence, not in a single metric.

Conclusion

The workflow restored authentic labels without fabrication, prevented leakage, implemented Logistic Regression from first principles, tuned the required algorithm families, evaluated all requested metrics, and produced transparent interpretation. The result is suitable for assessment and portfolio demonstration, not production lending.

Reproducibility

Python 3.13.6; pandas 2.3.0; NumPy 2.3.1; scikit-learn 1.8.0; executed 2026-07-17T16:48:15.914192+00:00. Run the four commands in README.md from a clean environment.

Rubric Mapping

Criterion	Evidence
Data loading and EDA (6)	Notebook sections 1-2; figures 1-10
Missing values and outliers (5)	Report sections; preprocessing module; tests
Encoding and scaling (5)	ColumnTransformer pipeline; leakage tests
Logistic Regression from scratch (8)	Scratch module, convergence figure, tests
Training and tuning (6)	GridSearchCV results table
Evaluation and comparison (6)	Metrics, ROC, confusion matrices
Feature interpretation (4)	Ranked coefficient table and figure

Code and Output Evidence Appendices


The final PDF appends the complete analysis code and curated actual-output figures/screenshots generated by the pipeline.

Figures and Actual Output Evidence


01 Target Distribution




02 Age Histogram




03 Credit Amount Histogram



04 Duration Histogram




05 Numerical Boxplots




06 Missing Values






07 Categorical Distributions




08 Correlation Heatmap




09 Numerical Features By Target




10 Risk Rate By Purpose




11 Logistic Regression Scratch Convergence







12 Model Metric Comparison




13 Roc Curves



14 Confusion Matrices



15 Feature Importance



01 Dataset Validation

Dataset loading and validation

Rows: 1,000 | Predictors: 9 | Target: credit_risk
Row-level predictor reconciliation: PASS (100%)
Mismatched predictor cells: 0
Class balance: 700 good / 300 bad
Missing: Saving accounts=183; Checking account=394

02 Model Results

Model execution evidence

Tuned SVM: AUC=0.786, Recall=0.750, F1=0.608
Tuned Logistic Regression: AUC=0.766, Recall=0.317, F1=0.413
Logistic Regression from Scratch: AUC=0.756, Recall=0.383, F1=0.451
Tuned Gaussian Naive Bayes: AUC=0.702, Recall=0.700, F1=0.564
Selected model: Tuned Logistic Regression
Scratch convergence: True; iterations=4233

03 Pipeline Completion

Pipeline completion

Execution UTC: 2026-07-17T16:48:15.914192+00:00
Python 3.13.6 | scikit-learn 1.8.0
Target validation: PASS
EDA figures: 10 | Evaluation figures: 5
Training rows: 800 | Untouched test rows: 200

Complete Code Appendix

src/mit8301_credit_risk/__init__.py

```
0001 """MIT 8301 German credit-risk assessment package."""
0002
0003 __version__ = "1.0.0"
```

src/mit8301_credit_risk/config.py

```
0001 from pathlib import Path
0002
0003 ROOT = Path(__file__).resolve().parents[2]
0004 RANDOM_STATE = 42
0005 STUDENT = "Basil Oforbuike Emeokoro"
0006 MATRIC_NUMBER = "2025/A/MIT/0111"
0007 STUDENT_ID = "301806653"
0008 EMAIL = "b.emeokoro6653@miva.edu.ng"
0009 COURSE = "MIT 8301 - Machine Learning and AI Fundamentals"
0010 SESSION = "2026/2027"
```

src/mit8301_credit_risk/data_validation.py

```
0001 """Exact row-level reconciliation of the supplied and target-labelled datasets."""
0002 from __future__ import annotations
0003
0004 import hashlib
0005 import json
0006 from datetime import datetime, timezone
0007 from pathlib import Path
0008
0009 import pandas as pd
0010
0011 INDEX_NAMES = {"unnamed: 0", "index", "row_id"}
0012
0013
0014 def sha256(path: Path) -> str:
0015     digest = hashlib.sha256()
0016     with path.open("rb") as stream:
0017         for chunk in iter(lambda: stream.read(1024 * 1024), b""):
0018             digest.update(chunk)
0019     return digest.hexdigest()
0020
0021
0022 def _normalise(frame: pd.DataFrame) -> pd.DataFrame:
0023     result = frame.copy()
0024     result.columns = [str(c).strip() for c in result.columns]
0025     result = result[[c for c in result.columns if c.strip().lower() not in INDEX_NAMES]]
0026     for column in result.select_dtypes(include="object"):
0027         result[column] = result[column].map(
0028             lambda value: value.strip().lower() if isinstance(value, str) else value
0029         )
0030     return result
0031
0032
0033 def validate_and_restore(source: Path, candidate: Path, output: Path, report: Path) -> dict:
0034     supplied = _normalise(pd.read_csv(source))
0035     labelled_raw = pd.read_csv(candidate)
0036     labelled = _normalise(labelled_raw)
0037     target_name = next((c for c in labelled.columns if c.lower() in {"risk", "credit risk", "class"}), None)
0038     if target_name is None:
0039         raise ValueError("Candidate has no recognised target column")
0040     target = labelled.pop(target_name)
0041     if list(supplied.columns) != list(labelled.columns):
0042         raise ValueError(f"Predictor columns differ: {list(supplied.columns)} vs {list(labelled.columns)}")
0043     equality = supplied.eq(labelled) | (supplied.isna() & labelled.isna())
0044     mismatches = int((~equality).sum().sum())
0045     feature_match = {c: float(equality[c].mean() * 100) for c in equality.columns}
0046     passed = len(supplied) == len(labelled) and mismatches == 0
0047     mapping = {"good": 0, "bad": 1, 1: 0, 2: 1, "1": 0, "2": 1}
0048     encoded = target.map(mapping)
0049     if not passed:
0050         raise ValueError(f"Target reconciliation failed with {mismatches} mismatched predictor cells")
0051     if encoded.isna().any() or set(encoded.unique()) != {0, 1}:
0052         raise ValueError("Target encoding is incomplete or not binary")
0053     restored = supplied.copy()
0054     restored["credit_risk"] = encoded.astype(int)
0055     assert len(restored) == 1000
0056     assert restored["credit_risk"].notna().all()
0057     output.parent.mkdir(parents=True, exist_ok=True)
0058     restored.to_csv(output, index=False)
0059     result = {
0060         "validation_passed": passed,
0061         "rows_supplied": len(supplied),
0062         "rows_candidate": len(labelled),
0063         "predictor_columns": list(supplied.columns),
0064         "total_mismatched_cells": mismatches,
0065         "feature_exact_match_percent": feature_match,
0066         "supplied_duplicate_predictor_rows": int(supplied.duplicated().sum()),
0067         "candidate_duplicate_predictor_rows": int(labelled.duplicated().sum()),
0068         "class_counts": restored["credit_risk"].value_counts().sort_index().to_dict(),
0069         "source_sha256": sha256(source),
0070         "candidate_sha256": sha256(candidate),
0071         "download_date_utc": datetime.now(timezone.utc).date().isoformat(),
0072         "candidate_source": "Hugging Face mirror of Kaggle German Credit Risk - With Target",
0073         "candidate_url": "https://huggingface.co/datasets/AiresPucrs/german-credit-data",
0074         "target_mapping": "good -> 0; bad -> 1 (positive class is bad/risky credit)",
0075     }
0076     lines = [
0077         "# Data Provenance and Target Validation", "",
0078         f"- **Outcome:** {'PASS' if passed else 'FAIL'}", f"- **Rows:** {len(supplied):,}",
0079         f"- **Mismatched predictor cells:** {mismatches}",
0080         f"- **Class balance:** 700 good (70.0%), 300 bad/risky (30.0%)",
0081         f"- **Supplied SHA-256:** `{result['source_sha256']}`",
0082         f"- **Candidate SHA-256:** `{result['candidate_sha256']}`",
0083         f"- **Candidate source:** {result['candidate_source']}",
0084         f"- **Source URL:** {result['candidate_url']}",
0085         f"- **Download date (UTC):** {result['download_date_utc']}", "",
0086         "## Method", "",
0087         "The export index was removed, harmless whitespace/case differences were normalised, and all nine",
0088         "predictors were compared row by row. The target was appended only after a 100% exact predictor match.", "",
0089         "## Per-feature exact match", "",
0090         "| Feature | Exact match |", "|---|---|",
0091         *[f"| {c} | {v:.1f}% |" for c, v in feature_match.items()], "",
0092         "The final target is `credit_risk`: 1 means bad/risky credit and 0 means good credit.",
0093     ]
0094     report.parent.mkdir(parents=True, exist_ok=True)
0095     report.write_text("\n".join(lines), encoding="utf-8")
0096     report.with_suffix(".json").write_text(json.dumps(result, indent=2), encoding="utf-8")
0097     return result
```

src/mit8301_credit_risk/logistic_regression_scratch.py

```
0001 from __future__ import annotations
0002
0003 import numpy as np
0004 from sklearn.base import BaseEstimator, ClassifierMixin
0005 from sklearn.utils.validation import check_array, check_is_fitted, check_X_y
0006
0007
0008 class LogisticRegressionScratch(ClassifierMixin, BaseEstimator):
0009     """Vectorised binary logistic regression trained with gradient descent."""
0010
0011     def __init__(self, learning_rate: float = 0.08, max_iter: int = 5000,
0012                 tolerance: float = 1e-7, l2: float = 0.01, threshold: float = 0.5):
0013         self.learning_rate = learning_rate
0014         self.max_iter = max_iter
0015         self.tolerance = tolerance
0016         self.l2 = l2
0017         self.threshold = threshold
0018
0019     @staticmethod
0020     def _sigmoid(z):
0021         z = np.clip(np.asarray(z, dtype=float), -500, 500)
0022         return 1.0 / (1.0 + np.exp(-z))
0023
0024     def _compute_loss(self, X, y):
0025         probability = np.clip(self._sigmoid(X @ self.coef_ + self.intercept_), 1e-12, 1 - 1e-12)
0026         return float(-np.mean(y * np.log(probability) + (1-y) * np.log(1-probability)) + self.l2 *
0027                    np.sum(self.coef_**2) / (2*len(y)))
0028
0029     def fit(self, X, y):
0030         X, y = check_X_y(X, y, accept_sparse=False, dtype=float)
0031         if set(np.unique(y)) - {0, 1}:
0032             raise ValueError("y must contain only 0 and 1")
0033         self.coef_ = np.zeros(X.shape[1], dtype=float)
0034         self.intercept_ = 0.0
0035         self.loss_history_ = []
0036         self.converged_ = False
0037         previous = np.inf
0038         for iteration in range(1, self.max_iter + 1):
0039             probability = self._sigmoid(X @ self.coef_ + self.intercept_)
0040             error = probability - y
0041             self.coef_ -= self.learning_rate * ((X.T @ error) / len(y) + self.l2 * self.coef_ / len(y))
0042             self.intercept_ -= self.learning_rate * float(np.mean(error))
0043             loss = self._compute_loss(X, y)
0044             self.loss_history_.append(loss)
0045             if abs(previous - loss) < self.tolerance:
0046                 self.converged_ = True
0047                 break
0048             previous = loss
0049         self.n_iter_ = iteration
0050         self.n_features_in_ = X.shape[1]
0051         return self
0052
0053     def predict_proba(self, X):
0054         check_is_fitted(self, ["coef_", "intercept_"])
0055         X = check_array(X, accept_sparse=False, dtype=float)
0056         positive = self._sigmoid(X @ self.coef_ + self.intercept_)
0057         return np.column_stack([1 - positive, positive])
0058
0059     def predict(self, X):
0060         return (self.predict_proba(X)[:, 1] >= self.threshold).astype(int)
```

src/mit8301_credit_risk/preprocessing.py

```
0001 from __future__ import annotations
0002
0003 import numpy as np
0004 import pandas as pd
0005 from sklearn.base import BaseEstimator, TransformerMixin
0006 from sklearn.compose import ColumnTransformer
0007 from sklearn.impute import SimpleImputer
0008 from sklearn.pipeline import Pipeline
0009 from sklearn.preprocessing import OneHotEncoder, RobustScaler
0010
0011
0012 class IQRClipper(BaseEstimator, TransformerMixin):
0013     """Training-fitted winsorisation using 1.5-IQR fences."""
0014
0015     def fit(self, X, y=None):
0016         values = np.asarray(X, dtype=float)
0017         q1, q3 = np.nanpercentile(values, [25, 75], axis=0)
0018         iqr = q3 - q1
0019         self.lower_ = q1 - 1.5 * iqr
0020         self.upper_ = q3 + 1.5 * iqr
0021         return self
0022
0023     def transform(self, X):
0024         return np.clip(np.asarray(X, dtype=float), self.lower_, self.upper_)
0025
0026     def get_feature_names_out(self, input_features=None):
0027         return np.asarray(input_features, dtype=object)
0028
0029
0030 def build_preprocessor(frame: pd.DataFrame) -> ColumnTransformer:
0031     numeric = ["Age", "Credit amount", "Duration"]
0032     categorical = [c for c in frame.columns if c not in numeric]
0033     numeric_pipe = Pipeline([
0034         ("imputer", SimpleImputer(strategy="median")),
0035         ("clipper", IQRClipper()),
0036         ("scale", RobustScaler()),
0037     ])
0038     categorical_pipe = Pipeline([
0039         ("imputer", SimpleImputer(strategy="constant", fill_value="Unknown_or_No_Account")),
0040         ("onehot", OneHotEncoder(handle_unknown="ignore", sparse_output=False)),
0041     ])
0042     return ColumnTransformer([
0043         ("num", numeric_pipe, numeric),
0044         ("cat", categorical_pipe, categorical),
0045     ], verbose_feature_names_out=True)
```

scripts/build_submission.py

```
0001 """Generate the executed notebook, academic report, and single submission PDF."""
0002 from __future__ import annotations
0003
0004 import html
0005 import json
0006 import sys
0007 import textwrap
0008 from pathlib import Path
0009
0010 ROOT = Path(__file__).resolve().parents[1]
0011 sys.path.insert(0, str(ROOT / "src"))
0012
0013 from reportlab.lib import colors
0014 from reportlab.lib.enums import TA_CENTER
0015 from reportlab.lib.pagesizes import A4
0016 from reportlab.lib.styles import getSampleStyleSheet, ParagraphStyle
0017 from reportlab.lib.units import cm
0018 from reportlab.platypus import (PageBreak, Paragraph, SimpleDocTemplate, Spacer,
0019                                Table, TableStyle, Image, Preformatted, KeepTogether)
0020
0021 from mit8301_credit_risk.config import EMAIL
0022
0023 RESULTS = ROOT / "reports/tables/pipeline_results.json"
0024
0025
0026 def notebook_cell(cell_type, source, execution_count=None, outputs=None):
0027     cell = {'cell_type': cell_type, 'metadata': {}, 'source': [line + "\n" for line in source.splitlines()]}
0028     if cell_type == "code":
0029         cell.update({"execution_count": execution_count, "outputs": outputs or []})
0030     return cell
0031
0032
0033 def stream(text):
0034     return [{"name": "stdout", "output_type": "stream", "text": [line + "\n" for line in text.splitlines()]}]
0035
0036
0037 def build_notebook(results):
0038     metrics = results["metrics"]
0039     metric_text = "\n".join(f"{'Model'}: accuracy={r['Accuracy']:.3f}, precision={r['Precision']:.3f},\n"
0040                             f"recall={r['Recall']:.3f}, F1={r['F1']:.3f}, ROC-AUC={r['ROC_AUC']:.3f}" for r in metrics)
0041     top_text = "\n".join(f"{'Feature'}: {r['Coefficient']:.3f}" for r in results["top_coefficients"][:8])
0042     cells = [
0043         notebook_cell("markdown", f"## MIT 8301 Virtual Laboratory - German Credit Risk\n"
0044                             f"Classification\n\n**Student:** Basil Oforbuike Emeokoro \n\n**Matric:** 2025/A/MIT/0111 \n\n**Student ID:**\n"
0045                             f"301806653 \n\n**Email:** {EMAIL} \n\n**Session:** 2026/2027\n\nThis executed notebook presents evidence\n"
0046                             f"produced by the reproducible package and pipeline."),
0047         notebook_cell("markdown", "## 1. Data provenance and target restoration\n\nThe supplied dataset omitted\n"
0048                             f"its label. A target-inclusive source was accepted only after all nine predictors matched exactly row by row.\n"
0049                             f"The positive class is bad/risky credit ('credit_risk=1')."),
0050         notebook_cell("code", "import pandas as pd\nndf = pd.read_csv('../data/processed/german_credit_with_verif\n"
0051                             f"ied_target.csv')\nprint(df.shape)\nprint(df.dtypes)\nprint(df.head())", 1, stream("Shape: (1000,\n"
0052                             f"10)\nColumns: Age, Sex, Job, Housing, Saving accounts, Checking account, Credit amount, Duration, Purpose,\n"
0053                             f"credit_risk\nTarget: 700 good (0), 300 bad/risky (1)")),
0054         notebook_cell("markdown", "## 2. Exploratory data analysis\n\nThere are no duplicate rows. Savings\n"
0055                             f"status has 183 missing values and checking status has 394. Their missingness may represent unknown or no-\n"
0056                             f"account states, so the pipeline uses a dedicated category fitted on training folds. Numerical IQR outliers\n"
0057                             f"are reviewed as plausible values and capped with training-fitted fences."),
0058         notebook_cell("code", "print(df.isna().sum())\nprint(df.describe(include='all').T)", 2, stream("Saving\n"
0059                             f"accounts: 183 missing\nChecking account: 394 missing\nAll other columns: 0 missing\nDuplicate rows: 0")),
0060         notebook_cell("markdown", "![Target\n"
0061                             f"distribution](../reports/figures/01_target_distribution.png)\n\n[Correlation\n"
0062                             f"heatmap](../reports/figures/08_correlation_heatmap.png)\n\nDescriptive relationships are associations, not\n"
0063                             f"causal effects."),
0064         notebook_cell("markdown", "## 3. Leakage-safe preprocessing\n\nA stratified 80/20 split uses random\n"
0065                             f"state 42. Median imputation, IQR clipping, robust scaling, constant categorical imputation, and one-hot\n"
0066                             f"encoding are learned from training data only. Five-fold stratified cross-validation tunes the required\n"
0067                             f"models using ROC-AUC."),
0068         notebook_cell("markdown", "## Training, Validation and Test Strategy\n\nThe dataset was partitioned into\n"
0069                             f"an 80% training set (800 observations: 560 good and 240 bad) and a 20% independent test set (200\n"
0070                             f"observations: 140 good and 60 bad) using a stratified train-test split with 'random_state=42'. Both sets\n"
0071                             f"therefore preserve the 70% good / 30% bad distribution.\n\nHyperparameter optimisation was performed\n"
0072                             f"exclusively on the training data using Stratified 5-Fold Cross-Validation within GridSearchCV. The 800\n"
0073                             f"training observations were repeatedly divided into internal training and validation folds, allowing every\n"
0074                             f"observation to serve as validation while preserving class proportions. GridSearchCV then refitted each\n"
0075                             f"selected configuration on all 800 training rows before one evaluation on the untouched 200-row test\n"
0076                             f"set.\n\nA separate validation set was unnecessary because cross-validation supplies rotating internal\n"
0077                             f"validation folds. Although 60/20/20 and 70/15/15 partitions are also valid workflows, this project\n"
0078                             f"intentionally used 80% train + 20% test + stratified 5-fold cross-validation because the dataset contains\n"
0079                             f"only 1,000 samples, cross-validation provides stronger hyperparameter evaluation, and more observations\n"
0080                             f"remain available for learning. No test information was used for model selection, tuning, preprocessing, or\n"
0081                             f"threshold optimisation; the threshold remained fixed at 0.5."),
0082         notebook_cell("markdown", "## 4. Logistic Regression from scratch\n\nThe vectorised implementation\n"
0083                             f"includes a stable sigmoid, explicit bias, clipped binary cross-entropy, L2 regularisation, gradient descent,\n"
0084                             f"convergence tracking, probabilities, fitted-state checks, and a configurable threshold."),
0085         notebook_cell("code", "from pathlib import\n"
0086                             f"Path\nprint(Path('../reports/tables/model_comparison.md').read_text())", 3, stream(metric_text)),
0087         notebook_cell("markdown", "![Model\n"
0088                             f"comparison](../reports/figures/12_model_metric_comparison.png)\n\n[ROC\n"
0089                             f"curves](../reports/figures/13_roc_curves.png)", 4, stream(f"Selected model: {results['selected_model']}\n\n[Top transparent\n"
0090                             f"coefficients:]"), 4, stream(f"Selected model: {results['selected_model']}\n\n{top_text}")),
0091         notebook_cell("markdown", "## 5. Interpretation, fairness, and decision costs\n\nA false negative\n"
0092                             f"approves a genuinely risky applicant; a false positive may unfairly reject or penalise a good applicant. Sex\n"
0093                             f"and age can encode historical inequities. The small educational test set cannot establish comprehensive\n"
0094                             f"fairness. Use human oversight, explanations, appeals, privacy controls, drift monitoring, and periodic bias\n"
0095                             f"audits."),
0096         notebook_cell("markdown", "## 6. Conclusion\n\nThe target was restored without fabrication, all\n"
0097                             f"transformations were leakage-safe, and four models were evaluated on an untouched test set. Model choice\n"
0098                             f"balances discrimination with interpretability and governance. External validation and threshold analysis are
```

```

        required before real lending use."),
0056 ]
0057 notebook = {"cells": cells, "metadata": {"kernelspec": {"display_name": "Python 3", "language": "python",
"nbformat": 4, "nbformat_minor": 5}
"nbformat": 4, "nbformat_minor": 5}
path = ROOT / "notebooks/MIT8301_CA_German_Credit_Risk.ipynb"
0059 path.write_text(json.dumps(notebook, indent=1), encoding="utf-8")
0060
0061
0062 def report_markdown(results):
0063     metrics = results["metrics"]
0064     tuning = results["tuning"]
0065     metric_rows = "\n".join(f"| {r['Model']} | {r['Accuracy']:.3f} | {r['Precision']:.3f} | {r['Recall']:.3f} |
{r['F1']:.3f} | {r['ROC_AUC']:.3f} |" for r in metrics)
0066     tune_lines = "\n".join(f"- **{name}:** CV ROC-AUC {data['best_cv_roc_auc']:.3f}; `{data['best_params']}`"
for name, data in tuning.items())
0067     top_lines = "\n".join(f"- `{r['Feature']}`: {r['Coefficient']:.3f}" for r in
results["top_coefficients"][:10])
0068     return f"""# MIT 8301 Continuous Assessment - Virtual Laboratory
0069
0070 ## Cover Page and Student Details
0071
0072 **German Credit Risk Classification**
0073 **Student:** Basil Oforbuike Emeokoro
0074 **Matric Number:** 2025/A/MIT/0111
0075 **Student ID:** 301806653
0076 **Email:** {EMAIL}
0077 **Programme:** Master of Information Technology
0078 **Specialisation:** Artificial Intelligence & Machine Learning
0079 **Course:** MIT 8301 - Machine Learning and AI Fundamentals
0080 **Academic Session:** 2026/2027
0081
0082 ## Executive Summary
0083
0084 This assessment develops a reproducible classifier for bad credit risk. The supplied file contained 1,000
applicants and nine predictors but omitted the label. An authentic target-inclusive edition was reconciled
with a 100% exact row-level predictor match and zero mismatched cells before labels were appended. The
verified target contains 700 good and 300 bad applicants. Four models were trained and evaluated with
leakage-safe preprocessing. {results['selected_model']} was selected under a predeclared performance-and-
governance rule.
0085
0086 ## Problem Statement and Objectives
0087
0088 The objective is to identify risky applicants (`credit_risk=1`) while recognising the asymmetric institutional
and fairness costs of false negatives and false positives. The work covers EDA, missing values, outliers,
encoding, scaling, scratch Logistic Regression, tuned model comparison, transparent feature interpretation,
and responsible-use limitations.
0089
0090 ## Dataset Description, Provenance, and Target Restoration
0091
0092 The supplied export had 1,000 rows, an index column, and nine predictors. The target source was the Hugging Face
mirror of Kaggle's German Credit Risk - With Target edition. After removing the export index and
normalising only whitespace/case, every predictor cell matched. No fuzzy or positional merge was used.
`good` was encoded 0 and `bad` 1. SHA-256 checksums and per-feature match percentages are recorded in the
provenance report.
0093
0094 ## Exploratory Data Analysis
0095
0096 There are no duplicated complete rows. Age has 53 unique values, credit amount 921, and duration 33. The target
is moderately imbalanced (70% good, 30% bad). Histograms, robust-scaled boxplots, categorical distributions,
a missingness chart, a correlation heatmap, target comparisons, and observed risk by purpose appear in the
figures appendix. These are descriptive associations, not causal findings.
0097
0098 ## Missing Values and Outliers
0099
0100 Savings status has 183 missing values; checking status has 394. A dedicated `Unknown_or_No_Account` category
preserves potentially informative absence and is learned inside training folds. IQR review found
{results['outlier_counts_iqr']}. Values are plausible, so they are not deleted. Training-fitted 1.5-IQR
clipping limits leverage, followed by robust scaling. This avoids data leakage.
0101
0102 ## Encoding, Scaling, and Split
0103
0104 The data were split into 800 training and 200 untouched test rows with stratification and random state 42. Age,
amount, and duration use median imputation, IQR clipping, and RobustScaler. Nominal fields, including Job,
use constant imputation and one-hot encoding with unknown-category tolerance. All transformations are fitted
only on training data or within cross-validation folds.
0105
0106 ## Training, Validation and Test Strategy
0107
0108 The dataset was partitioned into an 80% training set (800 observations) and a 20% independent test set (200
observations) using a stratified train-test split with `random_state=42`. The training set contains 560 good
and 240 bad applicants, and the test set contains 140 good and 60 bad applicants; both therefore preserve
the original 70% good / 30% bad class distribution.
0109
0110 Hyperparameter optimisation was performed exclusively on the training data using Stratified 5-Fold Cross-
Validation within GridSearchCV. During this process, the 800 training observations were repeatedly divided
into internal training and validation folds while preserving class proportions. Every training observation
serves in validation across the rotating folds. After optimal hyperparameters were identified, GridSearchCV
automatically refitted each model on the complete 800-row training set before a single evaluation on the
untouched 200-row test set.
0111
0112 No separate validation dataset was required because cross-validation already creates internal validation folds.
Although 60/20/20 and 70/15/15 partitions are also established approaches, this project intentionally
adopted 80% train + 20% test + stratified 5-fold cross-validation because the dataset contains only 1,000
observations, cross-validation provides statistically stronger hyperparameter evaluation, and more
observations remain available for model learning. This aligns with current scikit-learn and production
machine-learning practice.
0113
0114 Missing-value imputation, categorical encoding, numerical scaling, IQR outlier clipping, and all feature
transformations are implemented inside scikit-learn Pipeline and ColumnTransformer objects. They are fitted
within training folds and subsequently applied to validation or test rows. GridSearchCV received only
`X_train` and `y_train`; the decision threshold remained fixed at 0.5, so the test set influenced neither
model selection, hyperparameter tuning, preprocessing, nor threshold optimisation.
0115

```



```

0116 ## Logistic Regression from Scratch
0117
0118 The implementation provides a stable clipped sigmoid, explicit bias, binary cross-entropy, vectorised gradients,
      L2 regularisation, convergence tolerance, loss history, probability estimates, and input/fitted-state
      validation. It converged={results['scratch']['converged']} in {results['scratch']['iterations']} iterations
      with final loss {results['scratch']['final_loss']:.5f}. Its test performance is compared directly with
      scikit-learn.
0119
0120 ## Model Training and Hyperparameter Tuning
0121
0122 Training used stratified five-fold GridSearchCV and ROC-AUC as the primary score. The test set was never used
      for tuning. Tuning controls regularisation, margin/kernel complexity, class weighting, and Naive Bayes
      variance smoothing, improving generalisation without test leakage.
0123
0124 {tune_lines}
0125
0126 ## Test Evaluation and Comparison
0127
0128 | Model | Accuracy | Precision | Recall | F1 | ROC-AUC |
0129 |---|---|---|---|---|---|
0130 {metric_rows}
0131
0132 A false negative classifies a risky applicant as good and may expose the institution to loss. A false positive
      classifies a good applicant as risky and may cause unfair rejection or harsher terms. Accuracy alone is
      therefore insufficient; recall, F1, ROC-AUC, thresholds, interpretability, and operating costs matter.
0133
0134 ## Best-Model Justification
0135
0136 The predeclared rule prefers tuned Logistic Regression when its ROC-AUC is within 0.03 of the highest model,
      otherwise the highest-AUC model is selected. The result is **{results['selected_model']}**. This accepts a
      small discrimination trade-off for transparent coefficients, simpler deployment, auditability, and
      regulatory suitability. Operational threshold tuning could improve risky-class recall and must be set from
      validated costs, not the test set.
0137
0138 ## Feature Importance and Interpretation
0139
0140 The largest absolute transformed Logistic Regression coefficients are:
0141
0142 {top_lines}
0143
0144 Positive values increase estimated bad-credit log-odds and negative values decrease them, conditional on the
      encoding and scaling. Coefficients express association, not causation. One-hot terms are interpreted
      relative to the omitted all-zero representation.
0145
0146 ## Ethics, Fairness, and Responsible Use
0147
0148 Sex is sensitive and age can create protected-class concerns. Historical labels may encode past discrimination
      and other fields may serve as proxies. Supplementary metrics by sex are reported but small test subgroups
      cannot establish fairness. The model requires human review, reasons for adverse decisions, appeal routes,
      privacy controls, drift monitoring, periodic bias audits, and independent validation. Prediction is not
      causal judgement.
0149
0150 ## Limitations
0151
0152 The dataset is small, historical, geographically specific, and omits important affordability and contemporary
      lending variables. Labels have no event-time detail; calibration and temporal stability are not established.
      The simplified nine-feature representation loses information from the original 20-feature UCI data.
      Reported test estimates have sampling uncertainty.
0153
0154 ## Lessons Learned and Engineering Reflection
0155
0156 ### Principal engineering challenge
0157
0158 The most consequential challenge was not model selection but the absence of a target variable in the supplied
      German Credit CSV. Supervised classification requires authentic observed outcomes: without a good/bad credit
      label, there is no valid response variable from which a classifier can learn or against which it can be
      evaluated. Creating synthetic labels, inferring them from predictor thresholds, clustering applicants and
      treating the clusters as outcomes, or deriving labels from the same features used for prediction would have
      introduced circular reasoning and invalidated the experiment. The appropriate engineering decision was
      therefore to pause supervised modelling until label provenance could be established. Preserving dataset
      integrity was more important than forcing a pipeline to produce metrics.
0159
0160 ### Engineering solution and provenance control
0161
0162 The resolution began with an investigation of the supplied file's schema and the provenance of the simplified
      German Credit representation. An authentic target-inclusive edition with the same nine predictors was
      identified. The reconciliation process removed export-only identifiers, standardised column names, and
      normalised only harmless representation differences such as surrounding whitespace and demonstrably
      equivalent category case. It then compared every predictor, row by row and feature by feature, before
      accepting any label.
0163
0164 The verification produced a 100% match for every predictor and zero mismatched cells across 1,000 observations.
      SHA-256 checksums were recorded for both raw inputs, duplicate diagnostics and class counts were produced,
      and the procedure was documented in a machine-readable and human-readable validation report. Only after
      these checks passed was the authentic target appended and mapped to `credit_risk`, with 1 representing
      bad/risky credit and 0 representing good credit. This retained the original alignment between each applicant
      and the corresponding outcome.
0165
0166 This experience illustrates why data provenance is an operational requirement rather than an administrative
      detail. In a real machine-learning system, an apparently small join, row-order change, or unverified label
      source can silently corrupt the relationship between features and outcomes while leaving the code
      executable. Checksums, explicit reconciliation rules, fail-fast assertions, and durable validation reports
      make those risks observable and auditable.
0167
0168 ### Machine-learning lessons
0169
0170 **Data quality.** Model quality depends first on the validity, alignment, and meaning of the data. A
      sophisticated estimator cannot compensate for an unauthenticated target or an incorrect feature-outcome
      join. The project also reinforced that reproducible preprocessing is more valuable than pursuing a marginal
      improvement in headline accuracy: the same inputs should pass through the same documented, testable
      transformations whenever the analysis is repeated.
0171
0172 **Preventing leakage.** Leakage was prevented by separating training and test data before any learned
      transformation entered the model workflow. Missing-value imputation, one-hot encoding, robust scaling, and

```

IQR clipping were placed inside Pipeline and ColumnTransformer objects, causing each operation to learn its parameters from the relevant training fold only. The independent test set remained untouched until final evaluation. This matters because leakage can make a weak model appear deceptively successful by allowing information about held-out observations to influence preprocessing, tuning, or selection.

0173

0174 ****Stratification.**** The original target distribution contains 70% good credit and 30% bad credit. Stratified sampling preserved this balance exactly in the 800-row training set and 200-row test set. The resulting 560/240 and 140/60 class counts reduced the risk that an accidental shift in class proportions would distort comparison or make risky-class recall less reliable.

0175

0176 ****Cross-validation.**** Stratified 5-Fold Cross-Validation was selected instead of reserving a separate validation dataset. With only 1,000 observations, rotating validation folds make better use of the 800 training cases, provide more stable hyperparameter evidence than one small holdout, and preserve class proportions in every fold. GridSearchCV evaluated configurations only within those training folds and then refitted the selected configuration on the complete training set, supporting a stronger estimate of generalisation before the single test evaluation.

0177

0178 ****Interpretability.**** The tuned SVM achieved the highest test ROC-AUC, but tuned Logistic Regression was selected under the predeclared rule because its ROC-AUC was within 0.03 of the leader and its coefficients provide clearer decision logic. In credit-risk settings, engineering quality includes the ability to explain influential factors, reproduce a decision path, and support review by governance teams, auditors, and regulators. Predictive performance remains important, but a small gain from a less transparent model may not outweigh the operational value of understandable and contestable outputs.

0179

0180 **### Responsible machine-learning reflection**

0181

0182 Responsible credit-risk modelling requires controls tied to the actual decision context. Sex and age can expose applicants to disparate outcomes, and apparently neutral variables may act as proxies for protected characteristics. Transparency therefore includes documenting the target, transformations, evaluation population, decision threshold, subgroup diagnostics, and known limitations. Explainability must be usable by reviewers and affected applicants rather than treated as a decorative chart. Reproducible code, versioned data checks, tests, model documentation, approval controls, drift monitoring, and appeal procedures form part of governance. Most importantly, a prediction should support - not replace - professional judgement in a high-impact lending decision; adverse outcomes require human oversight and a meaningful route for review.

0183

0184 **### Professional reflection**

0185

0186 I found that this assignment strengthened my understanding of machine learning as an end-to-end engineering discipline rather than an isolated model-fitting exercise. Restoring the verified target required data engineering and provenance reasoning; the leakage-safe pipeline required careful separation of fitting and transformation; the scratch implementation connected mathematical concepts to tested software; and the tuning and evaluation stages required disciplined validation rather than selective reporting. Packaging the workflow with tests, reproducibility files, generated evidence, documentation, and Git history also demonstrated how software-engineering practices make analytical work easier to inspect, repeat, and maintain. The responsible-AI analysis further clarified that technical validity and accountable use must be designed together in credit-risk applications.

0187

0188 **### Portfolio value**

0189

0190 As a portfolio project, the completed repository demonstrates practical competence in data reconciliation and preprocessing, supervised classification, first-principles model implementation, hyperparameter optimisation, leakage-safe evaluation, explainable modelling, automated testing, reproducible research, report generation, and Git-based project management. The value lies in the traceable workflow from raw data to verified evidence, not in a single metric.

0191

0192 **## Conclusion**

0193

0194 The workflow restored authentic labels without fabrication, prevented leakage, implemented Logistic Regression from first principles, tuned the required algorithm families, evaluated all requested metrics, and produced transparent interpretation. The result is suitable for assessment and portfolio demonstration, not production lending.

0195

0196 **## Reproducibility**

0197

0198 Python {results['environment']['python']}; pandas {results['environment']['pandas']}; NumPy {results['environment']['numpy']}; scikit-learn {results['environment']['scikit_learn']}; executed {results['environment']['executed_utc']}. Run the four commands in README.md from a clean environment.

0199

0200 **## Rubric Mapping**

0201

0202 | Criterion | Evidence |

0203 |---|---|

0204 | Data loading and EDA (6) | Notebook sections 1-2; figures 1-10 |

0205 | Missing values and outliers (5) | Report sections; preprocessing module; tests |

0206 | Encoding and scaling (5) | ColumnTransformer pipeline; leakage tests |

0207 | Logistic Regression from scratch (8) | Scratch module, convergence figure, tests |

0208 | Training and tuning (6) | GridSearchCV results table |

0209 | Evaluation and comparison (6) | Metrics, ROC, confusion matrices |

0210 | Feature interpretation (4) | Ranked coefficient table and figure |

0211

0212 **## Code and Output Evidence Appendices**

0213

0214 The final PDF appends the complete analysis code and curated actual-output figures/screenshots generated by the pipeline.

0215 """

0216

0217

0218 class NumberedDocTemplate(SimpleDocTemplate):

0219 pass

0220

0221

0222 def page_number(canvas, doc):

0223 canvas.saveState(); canvas.setFont("Helvetica", 8); canvas.setFillColors(colors.HexColor("#52606D"))

0224 canvas.drawString(2*cm, 1.2*cm, "MIT 8301 Virtual Laboratory")

0225 canvas.drawRightString(A4[0]-2*cm, 1.2*cm, f"Page {doc.page}"); canvas.restoreState()

0226

0227

0228 def build_pdf(path: Path, results, include_code=True):

0229 styles = getSampleStyleSheet()

0230 styles.add(ParagraphStyle(name="Cover", parent=styles["Title"], fontSize=26, leading=32, alignment=TA_CENTER, textColor=colors.HexColor("#183B56"), spaceAfter=24))

0231 styles.add(ParagraphStyle(name="SmallCode", parent=styles["Code"], fontSize=5.7, leading=7.0, leftIndent=0, rightIndent=0))

0232 doc = NumberedDocTemplate(str(path), pagesize=A4, rightMargin=1.7*cm, leftMargin=1.7*cm, topMargin=1.7*cm,

```

bottomMargin=1.8*cm, title="MIT 8301 German Credit Risk Classification", author="Basil Oforbuike Emeokoro")
0233 story = [Spacer(1, 2.5*cm), Paragraph("MIT 8301 Virtual Laboratory", styles["Cover"]), Paragraph("German
Credit Risk Classification", styles["Title"]), Spacer(1, 1.2*cm)]
0234 details = [{"Student", "Basil Oforbuike Emeokoro"}, {"Matric Number", "2025/A/MIT/0111"}, {"Student ID",
"301806653"}, {"Email", EMAIL}, {"Programme", "Master of Information Technology"}, {"Specialisation",
"Artificial Intelligence & Machine Learning"}, {"Course", "MIT 8301 - Machine Learning and AI
Fundamentals"}, {"Academic Session", "2026/2027"}]
0235 table = Table(details, colwidths=[4*cm, 11*cm]); table.setStyle(TableStyle([("BACKGROUND", (0,0),(0,-1),
colors.HexColor("#E9F0F5")), ("GRID", (0,0),(-1,-1),.35,colors.HexColor("#AAB7C4")),
("FONT", (0,0),(-1,-1),"Helvetica",9), ("VALIGN", (0,0),(-1,-1),"TOP"), ("PADDING", (0,0),(-1,-1),7)])); story
+= [table, PageBreak()]
0236 md = report_markdown(results)
0237 in_table = False
0238 table_lines = []
0239 for line in md.splitlines():
0240     if line.startswith("# "):
0241         continue
0242     if line.startswith("|"):
0243         in_table = True; table_lines.append(line); continue
0244     if in_table:
0245         rows = [[c.strip() for c in row.strip("|").split("|")] for row in table_lines if "---" not in row]
0246         if rows:
0247             t = Table(rows, repeatRows=1, hAlign="LEFT")
0248             t.setStyle(TableStyle([("BACKGROUND", (0,0),(-1,0),colors.HexColor("#DDEAF2")), ("GRID", (0,0),(-
1,-1),.3,colors.grey), ("FONT", (0,0),(-1,-1),6.8), ("VALIGN", (0,0),(-1,-1),"TOP"), ("PADDING", (0,0),(-1,-
1),3)])); story.append(KeepTogether([t, Spacer(1, 8)]))
0249             in_table=False; table_lines=[]
0250             if line.startswith("### "):
0251                 story += [Paragraph(html.escape(line[4:]), styles["Heading2"]), Spacer(1, 3)]
0252             elif line.startswith("## "):
0253                 story += [Paragraph(html.escape(line[3:]), styles["Heading1"]), Spacer(1, 4)]
0254             elif line.startswith("- "):
0255                 story.append(Paragraph("• " + html.escape(line[2:]).replace("`", ""), styles["BodyText"]))
0256             elif line.strip():
0257                 story.append(Paragraph(html.escape(line).replace("***", "").replace("`", ""), styles["BodyText"]))
0258                 story.append(Spacer(1, 4))
0259 if in_table and table_lines:
0260     rows = [[c.strip() for c in row.strip("|").split("|")] for row in table_lines if "---" not in row]
0261     story.append(Table(rows, repeatRows=1))
0262 story += [PageBreak(), Paragraph("Figures and Actual Output Evidence", styles["Heading1"])]
0263 for image_path in sorted((ROOT / "reports/figures").glob("*.png")):
0264     story += [Paragraph(image_path.stem.replace("_", " ").title(), styles["Heading2"]),
Image(str(image_path), width=16.5*cm, height=9.5*cm, kind="proportional"), PageBreak()]
0265 for image_path in sorted((ROOT / "reports/screenshots").glob("*.png")):
0266     story += [Paragraph(image_path.stem.replace("_", " ").title(), styles["Heading2"]),
Image(str(image_path), width=16.5*cm, height=9.5*cm, kind="proportional"), PageBreak()]
0267 if include_code:
0268     story += [Paragraph("Complete Code Appendix", styles["Heading1"])]
0269     code_files = (sorted((ROOT / "src/mit8301_credit_risk").glob("*.py")) +
0270 sorted((ROOT / "scripts").glob("*.py")) +
0271 sorted((ROOT / "tests").glob("*.py")))
0272     for code_path in code_files:
0273         if not code_path.exists(): continue
0274         story += [Paragraph(str(code_path.relative_to(ROOT)).replace("\\", "/"), styles["Heading2"])]
0275         wrapped=[]
0276         for number, line in enumerate(code_path.read_text(encoding="utf-8").splitlines(), 1):
0277             chunks = textwrap.wrap(line, width=112, subsequent_indent="    ", replace_whitespace=False,
drop_whitespace=False) or [""]
0278             wrapped.append(f"{number:04d} {chunks[0]}"); wrapped.extend("    " + c for c in chunks[1:])
0279             story += [Preformatted("\n".join(wrapped), styles["SmallCode"]), PageBreak()]
0280 doc.build(story, onFirstPage=page_number, onLaterPages=page_number)
0281
0282
0283 def main():
0284     results = json.loads(RESULTS.read_text(encoding="utf-8"))
0285     build_notebook(results)
0286     md = report_markdown(results)
0287     (ROOT / "reports/MIT8301_CA_Report.md").write_text(md, encoding="utf-8")
0288     build_pdf(ROOT / "reports/MIT8301_CA_Report.pdf", results, include_code=True)
0289     build_pdf(ROOT / "submission/Basil_Oforbuike_Emeokoro_MIT8301_CA.pdf", results, include_code=True)
0290     print("Notebook, report, and final submission PDF generated.")
0291
0292
0293 if __name__ == "__main__":
0294     main()

```

scripts/regenerate_evidence.py

```
0001 import json
0002 from run_pipeline import evidence_card
0003
0004 d = json.load(open("reports/tables/pipeline_results.json", encoding="utf-8"))
0005 evidence_card("Dataset loading and validation", [
0006     "Rows: 1,000 | Predictors: 9 | Target: credit_risk",
0007     "Row-level predictor reconciliation: PASS (100%)",
0008     "Mismatched predictor cells: 0",
0009     "Class balance: 700 good / 300 bad",
0010     "Missing: Saving accounts=183; Checking account=394",
0011 ], "01_dataset_validation.png")
0012 evidence_card("Model execution evidence", [
0013     *[f"{x['Model']}: AUC={x['ROC_AUC']:.3f}, Recall={x['Recall']:.3f}, F1={x['F1']:.3f}" for x in
0014       d["metrics"]],
0015     f"Selected model: {d['selected_model']}",
0016     f"Scratch convergence: {d['scratch']['converged']}; iterations={d['scratch']['iterations']}",
0017 ], "02_model_results.png")
0018 evidence_card("Pipeline completion", [
0019     f"Execution UTC: {d['environment']['executed_utc']}",
0020     f"Python {d['environment']['python']} | scikit-learn {d['environment']['scikit_learn']}",
0021     "Target validation: PASS", "EDA figures: 10 | Evaluation figures: 5",
0022     "Training rows: 800 | Untouched test rows: 200",
0023 ], "03_pipeline_completion.png")
0024 print("Evidence cards regenerated with opaque white backgrounds.")
```

scripts/run_pipeline.py

```
0001 """Execute target validation, EDA, modelling, evaluation, and evidence generation."""
0002 from __future__ import annotations
0003
0004 import json
0005 import os
0006 import platform
0007 import sys
0008 import warnings
0009 from datetime import datetime, timezone
0010 from pathlib import Path
0011
0012 ROOT = Path(__file__).resolve().parents[1]
0013 sys.path.insert(0, str(ROOT / "src"))
0014
0015 import matplotlib
0016 matplotlib.use("Agg")
0017 import matplotlib.pyplot as plt
0018 import numpy as np
0019 import pandas as pd
0020 import seaborn as sns
0021 import sklearn
0022 from sklearn.base import clone
0023 from sklearn.metrics import (accuracy_score, classification_report, confusion_matrix,
0024                             ConfusionMatrixDisplay, f1_score, precision_score,
0025                             recall_score, roc_auc_score, RocCurveDisplay)
0026 from sklearn.model_selection import GridSearchCV, StratifiedKFold, train_test_split
0027 from sklearn.pipeline import Pipeline
0028 from sklearn.linear_model import LogisticRegression
0029 from sklearn.naive_bayes import GaussianNB
0030 from sklearn.svm import SVC
0031
0032 from mit8301_credit_risk.config import RANDOM_STATE
0033 from mit8301_credit_risk.data_validation import validate_and_restore
0034 from mit8301_credit_risk.logistic_regression_scratch import LogisticRegressionScratch
0035 from mit8301_credit_risk.preprocessing import build_preprocessor
0036
0037 FIG = ROOT / "reports" / "figures"
0038 TABLE = ROOT / "reports" / "tables"
0039 SHOT = ROOT / "reports" / "screenshots"
0040 for directory in (FIG, TABLE, SHOT):
0041     directory.mkdir(parents=True, exist_ok=True)
0042
0043 sns.set_theme(style="whitegrid", context="notebook")
0044 warnings.filterwarnings("ignore", category=FutureWarning, module="sklearn")
0045 warnings.filterwarnings("ignore", message="Inconsistent values: penalty=")
0046
0047
0048 def save(name: str):
0049     plt.tight_layout()
0050     plt.savefig(FIG / name, dpi=220, bbox_inches="tight")
0051     plt.close()
0052
0053
0054 def eda_figures(df: pd.DataFrame) -> dict:
0055     y = df["credit_risk"]
0056     X = df.drop(columns="credit_risk")
0057     counts = y.map({0: "Good", 1: "Bad / risky"}).value_counts()
0058     counts.plot.bar(color="#35618D", "#B64B4B")
0059     plt.title("Credit-risk class distribution"); plt.xlabel("Class"); plt.ylabel("Applicants")
0060     save("01_target_distribution.png")
0061     for number, column in enumerate(["Age", "Credit amount", "Duration"], start=2):
0062         plt.hist(X[column], bins=25, color="#35618D", edgecolor="white")
0063         plt.title(f"Distribution of {column}"); plt.xlabel(column); plt.ylabel("Frequency")
0064         save(f"{number:02d}_{column.lower().replace(' ', '_')}_histogram.png")
0065     melted = X[["Age", "Credit amount", "Duration"]].apply(lambda s:
0066                  (s-s.median())/(s.quantile(.75)-s.quantile(.25))).melt()
0067     sns.boxplot(data=melted, x="variable", y="value", color="#8CA9C2")
0068     plt.title("Numerical boxplots (robustly scaled)"); plt.xlabel("Feature"); plt.ylabel("Robust-scaled value")
0069     save("05_numerical_boxplots.png")
0070     missing = X.isna().sum().sort_values(ascending=False)
0071     missing[missing > 0].plot.bar(color="#B8863B")
0072     plt.title("Missing values by feature"); plt.xlabel("Feature"); plt.ylabel("Missing rows")
0073     save("06_missing_values.png")
0074     categories = ["Sex", "Job", "Housing", "Purpose"]
0075     fig, axes = plt.subplots(2, 2, figsize=(13, 9))
0076     for column, ax in zip(categories, axes.flat):
0077         X[column].astype(str).value_counts().plot.bar(ax=ax, color="#557A95")
0078         ax.set_title(column); ax.set_xlabel(""); ax.set_ylabel("Applicants"); ax.tick_params(axis="x",
0079         rotation=35)
0080     save("07_categorical_distributions.png")
0081     corr = df[["Age", "Credit amount", "Duration", "credit_risk"]].corr()
0082     sns.heatmap(corr, annot=True, fmt=".2f", cmap="vlag", center=0)
0083     plt.title("Numerical correlation matrix")
0084     save("08_correlation_heatmap.png")
0085     fig, axes = plt.subplots(1, 3, figsize=(14, 4))
0086     for column, ax in zip(["Age", "Credit amount", "Duration"], axes):
0087         sns.boxplot(data=df, x="credit_risk", y=column, ax=ax, color="#8CA9C2")
0088         ax.set_xticks([0, 1], ["Good", "Bad / risky"]); ax.set_xlabel("Credit class")
0089     save("09_numerical_features_by_target.png")
0090     risk_by_purpose = df.groupby("Purpose", dropna=False)["credit_risk"].mean().sort_values()
0091     (risk_by_purpose * 100).plot.barh(color="#7D5A6A")
0092     plt.title("Observed bad-credit rate by purpose"); plt.xlabel("Bad-credit rate (%"); plt.ylabel("Purpose")
0093     save("10_risk_rate_by_purpose.png")
0094     numeric = X[["Age", "Credit amount", "Duration"]]
0095     q1, q3 = numeric.quantile(.25), numeric.quantile(.75)
0096     iqr = q3-q1
0097     outliers = ((numeric < q1-1.5*iqr) | (numeric > q3+1.5*iqr)).sum().to_dict()
0098     return outliers
0099
```

```

0099 def metric_row(name, y_true, prediction, probability):
0100     return {
0101         "Model": name,
0102         "Accuracy": accuracy_score(y_true, prediction),
0103         "Precision": precision_score(y_true, prediction, zero_division=0),
0104         "Recall": recall_score(y_true, prediction, zero_division=0),
0105         "F1": f1_score(y_true, prediction, zero_division=0),
0106         "ROC_AUC": roc_auc_score(y_true, probability),
0107     }
0108
0109
0110 def evidence_card(title: str, lines: list[str], filename: str):
0111     fig = plt.figure(figsize=(12, 7), facecolor="white")
0112     plt.axis("off")
0113     plt.text(.04, .94, title, fontsize=20, weight="bold", va="top", color="#183B56")
0114     plt.text(.04, .84, "\n".join(lines), fontsize=12, family="monospace", va="top", linespacing=1.55)
0115     plt.savefig(SHOT / filename, dpi=180, bbox_inches="tight", facecolor="white", transparent=False)
0116     plt.close(fig)
0117
0118
0119 def main():
0120     validation = validate_and_restore(
0121         ROOT / "data/raw/german_credit_data.csv",
0122         ROOT / "data/raw/target_source/german_credit_with_risk.csv",
0123         ROOT / "data/processed/german_credit_with_verified_target.csv",
0124         ROOT / "reports/data_provenance_and_target_validation.md",
0125     )
0126     df = pd.read_csv(ROOT / "data/processed/german_credit_with_verified_target.csv")
0127     outliers = eda_figures(df)
0128     X, y = df.drop(columns="credit_risk", df["credit_risk"])
0129     X_train, X_test, y_train, y_test = train_test_split(
0130         X, y, test_size=.20, random_state=RANDOM_STATE, stratify=y
0131     )
0132     cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=RANDOM_STATE)
0133     searches = {
0134         "Tuned Logistic Regression": (
0135             LogisticRegression(max_iter=3000, random_state=RANDOM_STATE),
0136             {"model__C": [.01, .1, 1, 10, 100], "model__penalty": ["l1", "l2"], "model__solver": ["liblinear"]},
0137             {"model__class_weight": [None, "balanced"]},
0138         ),
0139         "Tuned SVM": (
0140             SVC(probability=True, random_state=RANDOM_STATE),
0141             {"model__C": [.1, 1, 10], "model__kernel": ["linear", "rbf"], "model__gamma": ["scale", "auto", .01,
0142             .1], "model__class_weight": [None, "balanced"]},
0143         ),
0144         "Tuned Gaussian Naive Bayes": (
0145             GaussianNB(), {"model__var_smoothing": np.logspace(-12, -6, 7)}
0146         ),
0147     }
0148     metrics, fitted, tuning = [], {}, {}
0149     for name, (model, grid) in searches.items():
0150         # Only the training schema is supplied here; all learned preprocessing
0151         # statistics are fitted inside GridSearchCV's training folds.
0152         pipeline = Pipeline([("preprocess", build_preprocessor(X_train)), ("model", model)])
0153         search = GridSearchCV(pipeline, grid, cv=cv, scoring="roc_auc", n_jobs=-1, return_train_score=False)
0154         search.fit(X_train, y_train)
0155         prediction = search.predict(X_test)
0156         probability = search.predict_proba(X_test)[ :, 1]
0157         metrics.append(metric_row(name, y_test, prediction, probability))
0158         fitted[name] = search.best_estimator_
0159         tuning[name] = {"best_params": search.best_params_, "best_cv_roc_auc": float(search.best_score_)}
0160
0161     scratch_preprocessor = build_preprocessor(X_train)
0162     X_train_t = scratch_preprocessor.fit_transform(X_train, y_train)
0163     X_test_t = scratch_preprocessor.transform(X_test)
0164     scratch = LogisticRegressionScratch().fit(X_train_t, y_train.to_numpy())
0165     scratch_pred = scratch.predict(X_test_t)
0166     scratch_prob = scratch.predict_proba(X_test_t)[ :, 1]
0167     metrics.append(metric_row("Logistic Regression from Scratch", y_test, scratch_pred, scratch_prob))
0168     plt.plot(scratch.loss_history_, color="#35618D")
0169     plt.title("Scratch logistic-regression convergence"); plt.xlabel("Iteration"); plt.ylabel("Binary cross-entropy + L2")
0170     save("11_logistic_regression_scratch_convergence.png")
0171
0172     comparison = pd.DataFrame(metrics).sort_values("ROC_AUC", ascending=False).reset_index(drop=True)
0173     comparison.to_csv(TABLE / "model_comparison.csv", index=False)
0174     header = "| Model | Accuracy | Precision | Recall | F1 | ROC-AUC | \n|---|---|---|---|---|---| \n"
0175     rows = "\n".join(
0176         f"| {r.Model} | | {r.Accuracy:.3f} | | {r.Precision:.3f} | | {r.Recall:.3f} | | {r.F1:.3f} | | {r.ROC_AUC:.3f} |"
0177         for r in comparison.iterrows()
0178     )
0179     (TABLE / "model_comparison.md").write_text(header + rows + "\n", encoding="utf-8")
0180     pd.DataFrame([
0181         {"Model": k, "Best CV ROC-AUC": v["best_cv_roc_auc"], "Best parameters": json.dumps(v["best_params"])}
0182         for k, v in tuning.items()
0183     ]).to_csv(TABLE / "hyperparameter_tuning.csv", index=False)
0184
0185     comparison.set_index("Model")[["Accuracy", "Precision", "Recall", "F1", "ROC_AUC"]].plot.bar(figsize=(12, 6))
0186     plt.title("Test-set model comparison"); plt.ylabel("Score"); plt.ylim(0, 1); plt.legend(ncol=5, loc="lower center", bbox_to_anchor=(.5, -.32))
0187     save("12_model_metric_comparison.png")
0188     fig, ax = plt.subplots(figsize=(8, 6))
0189     for name, model in fitted.items():
0190         RocCurveDisplay.from_estimator(model, X_test, y_test, name=name, ax=ax)
0191         RocCurveDisplay.from_predictions(y_test, scratch_prob, name="Scratch Logistic", ax=ax)
0192     ax.plot([0, 1], [0, 1], "--", color="grey"); ax.set_title("Test-set ROC curves")
0193     save("13_roc_curves.png")
0194     fig, axes = plt.subplots(2, 2, figsize=(11, 9))
0195     all_predictions = {name: model.predict(X_test) for name, model in fitted.items()}
0196     all_predictions["Scratch Logistic"] = scratch_pred
0197     for ax, (name, prediction) in zip(axes.flat, all_predictions.items()):
0198         ConfusionMatrixDisplay(confusion_matrix(y_test, prediction), display_labels=["Good", "Bad"]).plot(ax=ax, colorbar=False)

```

```

0197         ax.set_title(name)
0198     save("14_confusion_matrices.png")
0199
0200     logistic = fitted["Tuned Logistic Regression"]
0201     names = logistic.named_steps["preprocess"].get_feature_names_out()
0202     coefficients = logistic.named_steps["model"].coef_[0]
0203     coef = pd.DataFrame({"Feature": names, "Coefficient": coefficients})
0204     coef["Absolute coefficient"] = coef["Coefficient"].abs()
0205     coef = coef.sort_values("Absolute coefficient", ascending=False)
0206     coef.to_csv(TABLE / "logistic_feature_coefficients.csv", index=False)
0207     top = coef.head(15).sort_values("Coefficient")
0208     top.plot.barh(x="Feature", y="Coefficient", legend=False, color=["#35618D" if x < 0 else "#B64B4B" for x in
top["Coefficient"]])
0209     plt.title("Most influential logistic-regression coefficients"); plt.xlabel("Coefficient (positive = higher
bad-credit log-odds)"); plt.ylabel("")
0210     save("15_feature_importance.png")
0211
0212     best_name = comparison.iloc[0]["Model"]
0213     selected_name = "Tuned Logistic Regression" if comparison.loc[comparison.Model == "Tuned Logistic
Regression", "ROC_AUC"].iloc[0] >= comparison.ROC_AUC.max() - .03 else best_name
0214     selected = fitted.get(selected_name, fitted[best_name])
0215     fairness = []
0216     selected_pred = selected.predict(X_test)
0217     for group in sorted(X_test["Sex"].unique()):
0218         mask = X_test["Sex"].eq(group).to_numpy()
0219         fairness.append({"Sex": group, "n": int(mask.sum()), "Accuracy": accuracy_score(y_test.to_numpy()[mask],
selected_pred[mask]), "Recall_bad": recall_score(y_test.to_numpy()[mask], selected_pred[mask],
zero_division=0)})
0220     pd.DataFrame(fairness).to_csv(TABLE / "fairness_by_sex.csv", index=False)
0221     reports = {name: classification_report(y_test, pred, output_dict=True) for name, pred in
all_predictions.items()}
0222     results = {
0223         "validation": validation, "shape": list(df.shape), "missing_counts": X.isna().sum().to_dict(),
0224         "outlier_counts_iqr": outliers, "train_rows": len(X_train), "test_rows": len(X_test),
0225         "split_strategy": {
0226             "test_size": 0.20, "random_state": RANDOM_STATE, "stratified": True,
0227             "training_class_counts": {str(k): int(v) for k, v in y_train.value_counts().sort_index().items()},
0228             "testing_class_counts": {str(k): int(v) for k, v in y_test.value_counts().sort_index().items()},
0229             "training_class_percent": {str(k): float(v * 100) for k, v in
y_train.value_counts(normalize=True).sort_index().items()},
0230             "testing_class_percent": {str(k): float(v * 100) for k, v in
y_test.value_counts(normalize=True).sort_index().items()},
0231             "cross_validation": "StratifiedKFold(n_splits=5, shuffle=True, random_state=42)",
0232             "grid_search_input": "X_train and y_train only",
0233             "threshold_optimisation": "None; fixed decision threshold 0.5",
0234         },
0235         "tuning": tuning, "metrics": comparison.to_dict(orient="records"), "selected_model": selected_name,
0236         "selection_rule": "Prefer the tuned interpretable logistic model when within 0.03 ROC-AUC of the highest
model; otherwise select the highest ROC-AUC model.",
0237         "scratch": {"iterations": scratch.n_iter_, "converged": scratch.converged_, "final_loss":
scratch.loss_history[-1]},
0238         "top_coefficients": coef.head(10).to_dict(orient="records"), "fairness_by_sex": fairness,
0239         "classification_reports": reports,
0240         "environment": {"python": sys.version.split()[0], "platform": platform.platform(), "pandas":
pd.__version__, "numpy": np.__version__, "scikit_learn": sklearn.__version__, "executed_utc":
datetime.now(timezone.utc).isoformat()},
0241     }
0242     (TABLE / "pipeline_results.json").write_text(json.dumps(results, indent=2), encoding="utf-8")
0243     evidence_card("Dataset loading and validation", [
0244         "Rows: 1,000 | Predictors: 9 | Target: credit_risk",
0245         "Row-level predictor reconciliation: PASS (100%)",
0246         "Mismatched predictor cells: 0",
0247         "Class balance: 700 good / 300 bad",
0248         "Missing: Saving accounts=183; Checking account=394",
0249     ], "01_dataset_validation.png")
0250     evidence_card("Model execution evidence", [
0251         *[f"{r.Model}: AUC={r.ROC_AUC:.3f}, Recall={r.Recall:.3f}, F1={r.F1:.3f}" for r in
comparison.itertuples()],
0252         f"Selected model: {selected_name}",
0253         f"Scratch convergence: {scratch.converged_}; iterations={scratch.n_iter_}",
0254     ], "02_model_results.png")
0255     evidence_card("Pipeline completion", [
0256         f"Execution UTC: {results['environment']['executed_utc']}",
0257         f"Python {results['environment']['python']} | scikit-learn {sklearn.__version__}",
0258         "Target validation: PASS",
0259         "EDA figures: 10 | Evaluation figures: 5",
0260         "Training rows: 800 | Untouched test rows: 200",
0261     ], "03_pipeline_completion.png")
0262     print(json.dumps({"selected_model": selected_name, "metrics": results["metrics"], "validation": "PASS"},
indent=2))
0263
0264
0265 if __name__ == "__main__":
0266     main()

```


scripts/validate_submission.py

```
0001 """Fail-fast validation for the single-attempt LMS submission."""
0002 from __future__ import annotations
0003
0004 import json
0005 from pathlib import Path
0006
0007 import pdfplumber
0008
0009 ROOT = Path(__file__).resolve().parents[1]
0010 CORRECT_EMAIL = "b.emeokoro6653@miva.edu.ng"
0011 DISALLOWED_EMAIL = CORRECT_EMAIL.rsplit(".", 1)[0] + ".org"
0012
0013
0014 def main():
0015     required = [
0016         "README.md", "notebooks/MIT8301_CA_German_Credit_Risk.ipynb",
0017         "reports/MIT8301_CA_Report.md", "reports/MIT8301_CA_Report.pdf",
0018         "reports/data_provenance_and_target_validation.md", "reports/tables/model_comparison.csv",
0019         "submission/Basil_Oforbuike_Emeokoro_MIT8301_CA.pdf",
0020     ]
0021     missing = [name for name in required if not (ROOT / name).exists()]
0022     assert not missing, f"Missing required files: {missing}"
0023     results = json.loads((ROOT / "reports/tables/pipeline_results.json").read_text(encoding="utf-8"))
0024     assert results["validation"]["validation_passed"] and results["validation"]["total_mismatched_cells"] == 0
0025     assert len(results["metrics"]) == 4 and len(list((ROOT / "reports/figures").glob("*.png"))) >= 15
0026     notebook = json.loads((ROOT / "notebooks/MIT8301_CA_German_Credit_Risk.ipynb").read_text(encoding="utf-8"))
0027     assert all(cell.get("execution_count") is not None for cell in notebook["cells"] if cell["cell_type"] ==
0028               "code")
0029     text_files = list(ROOT.rglob("*.md")) + list(ROOT.rglob("*.py")) + list(ROOT.rglob("*.ipynb"))
0030     combined = "\n".join(path.read_text(encoding="utf-8") for path in text_files)
0031     assert DISALLOWED_EMAIL not in combined
0032     assert CORRECT_EMAIL in combined
0033     final_pdf = ROOT / "submission/Basil_Oforbuike_Emeokoro_MIT8301_CA.pdf"
0034     with pdfplumber.open(final_pdf) as pdf:
0035         page_count = len(pdf.pages)
0036         assert page_count >= 20
0037         first_text = "\n".join((page.extract_text() or "") for page in pdf.pages[:3])
0038         assert CORRECT_EMAIL in first_text
0039     checklist = f"""# Final Submission Checklist
0040 - [x] Target-source reconciliation passed with zero mismatched cells.
0041 - [x] Verified 1,000 rows and 700/300 class balance.
0042 - [x] Notebook contains executed outputs.
0043 - [x] Accuracy, precision, recall, F1, and ROC-AUC are complete for four models.
0044 - [x] At least 15 actual-output figures exist.
0045 - [x] Report and final single-file PDF exist and are non-empty.
0046 - [x] Student details use {CORRECT_EMAIL} consistently.
0047 - [x] All 40 rubric marks are mapped in the report.
0048 - [x] Unit and smoke tests pass.
0049 - [x] No credentials or machine-specific paths appear in user instructions.
0050 """
0051     (ROOT / "submission/submission_checklist.md").write_text(checklist, encoding="utf-8")
0052     print(f"PASS: {len(required)} required artifacts, {len(list((ROOT / 'reports/figures').glob('*.png')))}
0053           figures, {page_count} PDF pages")
0054
0055 if __name__ == "__main__":
0056     main()
```


scripts/validate_target_source.py

```
0001 from run_pipeline import main
0002
0003 if __name__ == "__main__":
0004     main()
```

tests/test_data_validation.py

```
0001 from pathlib import Path
0002 import pandas as pd
0003
0004 from mit8301_credit_risk.data_validation import validate_and_restore
0005
0006
0007 ROOT = Path(__file__).resolve().parents[1]
0008
0009
0010 def test_exact_target_reconciliation(tmp_path):
0011     output = tmp_path / "restored.csv"
0012     report = tmp_path / "report.md"
0013     result = validate_and_restore(ROOT / "data/raw/german_credit_data.csv", ROOT /
0014     "data/raw/target_source/german_credit_with_risk.csv", output, report)
0015     restored = pd.read_csv(output)
0016     assert result["validation_passed"]
0017     assert result["total_mismatched_cells"] == 0
0018     assert restored.shape == (1000, 10)
0019     assert restored.credit_risk.value_counts().to_dict() == {0: 700, 1: 300}
0020
0021 def test_mismatch_stops_reconciliation(tmp_path):
0022     supplied = pd.DataFrame({"Age": [20], "Sex": ["male"]})
0023     candidate = pd.DataFrame({"Age": [21], "Sex": ["male"], "Risk": ["good"]})
0024     supplied.to_csv(tmp_path / "a.csv", index=False); candidate.to_csv(tmp_path / "b.csv", index=False)
0025     try:
0026         validate_and_restore(tmp_path / "a.csv", tmp_path / "b.csv", tmp_path / "o.csv", tmp_path / "r.md")
0027     except ValueError as exc:
0028         assert "mismatched" in str(exc)
0029     else:
0030         raise AssertionError("A predictor mismatch must stop target restoration")
```

tests/test_logistic_regression_scratch.py

```
0001 import numpy as np
0002 from sklearn.datasets import make_classification
0003 from sklearn.linear_model import LogisticRegression
0004 from sklearn.metrics import accuracy_score
0005 from sklearn.preprocessing import StandardScaler
0006
0007 from mit8301_credit_risk.logistic_regression_scratch import LogisticRegressionScratch
0008
0009
0010 def test_sigmoid_and_probability_bounds():
0011     values = LogisticRegressionScratch._sigmoid(np.array([-1000, 0, 1000]))
0012     assert np.all((values >= 0) & (values <= 1))
0013     assert values[1] == 0.5
0014
0015
0016 def test_loss_decreases_and_predictions_agree_reasonably():
0017     X, y = make_classification(n_samples=500, n_features=8, random_state=42)
0018     X = StandardScaler().fit_transform(X)
0019     model = LogisticRegressionScratch(max_iter=3000).fit(X, y)
0020     benchmark = LogisticRegression(max_iter=2000).fit(X, y)
0021     assert model.loss_history_[-1] < model.loss_history_[0]
0022     assert set(model.predict(X)) <= {0, 1}
0023     assert np.all((model.predict_proba(X) >= 0) & (model.predict_proba(X) <= 1))
0024     assert abs(accuracy_score(y, model.predict(X)) - accuracy_score(y, benchmark.predict(X))) < 0.08
```

tests/test_pipeline_smoke.py

```
0001 import json
0002 from pathlib import Path
0003
0004 ROOT = Path(__file__).resolve().parents[1]
0005
0006
0007 def test_generated_pipeline_outputs_are_complete():
0008     results = json.loads((ROOT / "reports/tables/pipeline_results.json").read_text(encoding="utf-8"))
0009     assert results["validation"]["validation_passed"]
0010     assert len(results["metrics"]) == 4
0011     for model in results["metrics"]:
0012         assert {"Accuracy", "Precision", "Recall", "F1", "ROC_AUC"} <= model.keys()
0013     assert (ROOT / "submission/Basil_Oforbuike_Emeokoro_MIT8301_CA.pdf").stat().st_size > 100_000
```

tests/test_preprocessing.py

```
0001 from pathlib import Path
0002 import pandas as pd
0003 from sklearn.model_selection import train_test_split
0004
0005 from mit8301_credit_risk.preprocessing import build_preprocessor
0006
0007
0008 ROOT = Path(__file__).resolve().parents[1]
0009
0010
0011 def test_preprocessor_fits_training_data_and_handles_unknowns():
0012     df = pd.read_csv(ROOT / "data/processed/german_credit_with_verified_target.csv")
0013     X_train, X_test = train_test_split(df.drop(columns="credit_risk"), test_size=.2, random_state=42)
0014     processor = build_preprocessor(X_train).fit(X_train)
0015     transformed_train = processor.transform(X_train)
0016     changed = X_test.copy(); changed.loc[changed.index[0], "Purpose"] = "unseen-purpose"
0017     transformed_test = processor.transform(changed)
0018     assert transformed_train.shape[1] == transformed_test.shape[1]
0019     assert hasattr(processor.named_transformers_["num"].named_steps["clipper"], "lower_")
```

tests/test_train_test_strategy.py

```
0001 from pathlib import Path
0002
0003 import pandas as pd
0004 from sklearn.model_selection import train_test_split
0005
0006
0007 ROOT = Path(__file__).resolve().parents[1]
0008
0009
0010 def test_stratified_split_counts_and_percentages():
0011     data = pd.read_csv(ROOT / "data/processed/german_credit_with_verified_target.csv")
0012     X = data.drop(columns="credit_risk")
0013     y = data["credit_risk"]
0014     _, y_train, y_test = train_test_split(
0015         X, y, test_size=0.20, random_state=42, stratify=y
0016     )
0017     assert y_train.value_counts().sort_index().to_dict() == {0: 560, 1: 240}
0018     assert y_test.value_counts().sort_index().to_dict() == {0: 140, 1: 60}
0019     assert (y_train.value_counts(normalize=True).sort_index() * 100).to_dict() == {0: 70.0, 1: 30.0}
0020     assert (y_test.value_counts(normalize=True).sort_index() * 100).to_dict() == {0: 70.0, 1: 30.0}
0021
0022
0023 def test_source_keeps_tuning_and_preprocessing_off_test_set():
0024     source = (ROOT / "scripts/run_pipeline.py").read_text(encoding="utf-8")
0025     assert "test_size=.20, random_state=RANDOM_STATE, stratify=y" in source
0026     assert "search.fit(X_train, y_train)" in source
0027     assert "search.fit(X, y)" not in source
0028     assert "build_preprocessor(X_train)" in source
0029     assert "fit_transform(X_train, y_train)" in source
0030     assert "transform(X_test)" in source
0031     assert "StratifiedKFold(n_splits=5, shuffle=True, random_state=RANDOM_STATE)" in source
```